

Experiment: Data pipelines using Azure Blob storage, Azure SQL Database, and Azure Data Factory

In this experiment, you create a data factory by using the Azure Data Factory user interface (UI). The pipeline in this data factory copies data from Azure Blob storage to a database in Azure SQL Database. The configuration pattern in this experiment applies to copying from a file-based data store to a relational data store. For a list of data stores supported as sources and sinks, see the [supported data stores](#) table.

Note

Prerequisites

- **Azure subscription.** If you don't have an Azure subscription, create a [free Azure account](#) before you begin.
- **Azure storage account.** You use Blob storage as a *source* data store. If you don't have a storage account, see [Create an Azure storage account](#) for steps to create one.
- **Azure SQL Database.** You use the database as a *sink* data store. If you don't have a database in Azure SQL Database, see the [Create a database in Azure SQL Database](#) for steps to create one.

How to do it...

The steps for this experiment are as follows:

1. Execute the following commands to create the employee container in our Azure storage account:

```
$storageaccountname="adestoragepowershellxx"
$containername="input"
$resourcegroup="AzureDataEngineeringxx"
#Get the Azure Storage account context
$storagecontext = (Get-AzStorageAccount -ResourceGroupName
$resourcegroup -Name $storageaccountname).Context;
#Create a new container
New-AzStorageContainer -Name $containername -Context
$storagecontext
```

Container creation is usually very quick.

Note: Substitute for the student number for xx, (student1 would be adestoragepowershell01, student12 would be adestoragepowershell12). Similarly as we did earlier for the AzureDataEngineeringxx ResourceGroupName (student1 would be AzureDataEngineering01, student12 would be AzureDataEngineering12).

```
PS C:\projects\AzureDataEngineering> #Create a new container
>> New-AzStorageContainer -Name $containername -Context $storagecontext

Storage Account Name: adestoragepowershell

Name                PublicAccess      LastModified      IsDeleted  VersionId
----                -
input              Off               5/19/2022 12:35:22 PM +00:00
```

Create a blob and a SQL table

Create a source blob

1. Launch Notepad. Copy the following text, and save it as an **emp.txt** file on your disk:

```
FirstName,LastName
John,Doe
Jane,Doe
```

2. We can use CLI or the [Azure Storage Explorer](#) to upload this file to our input container created earlier in the experience.

Set-AzStorageBlobContent -File "C:\projects\AzureDataEngineering\repos\azure-data-engineering\labs\emp.txt" -Context \$storagecontext -Blob emp.txt -Container \$containername

```
PS C:\projects\AzureDataEngineering> Set-AzStorageBlobContent -File "C:\projects\AzureDataEngineering\repos\azure-data-engineering\labs\emp.txt" -Context $storagecontext -Blob emp.txt -Container $containername

AccountName: adestoragepowershell, ContainerName: input

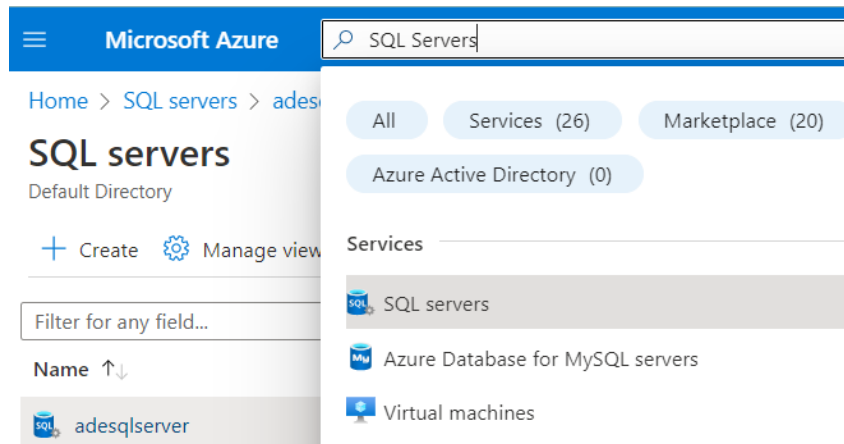
Name                BlobType  Length  ContentType      LastModified      Access
----                -
emp.txt             BlockBlob 38      application/octet-stream 2022-05-19 12:36:20Z Hot
```

Create a sink SQL table

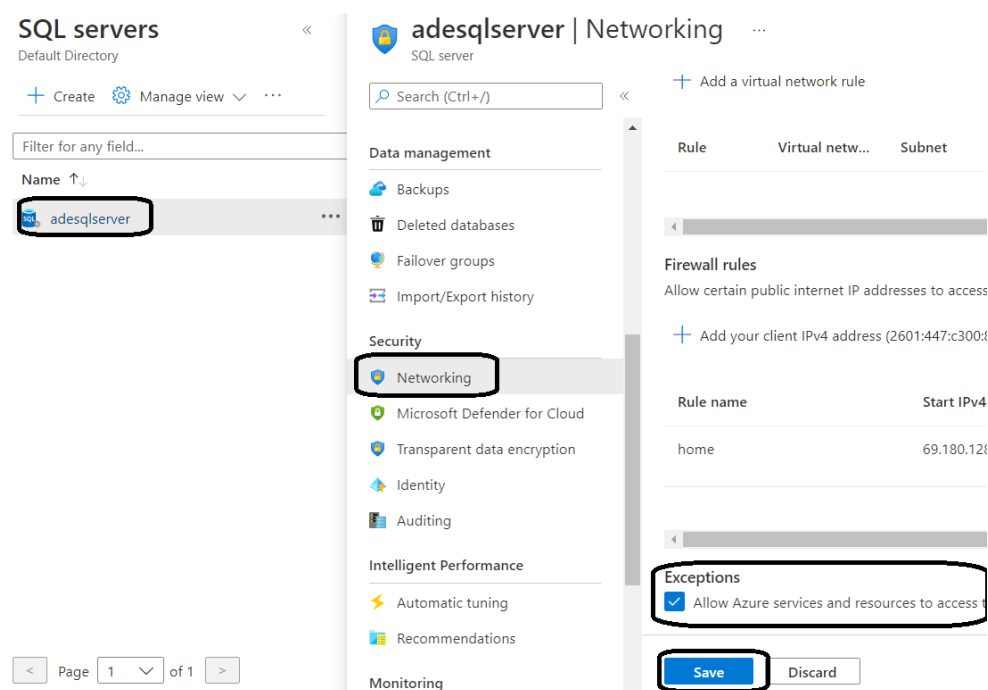
1. The SQL Server PowerShell module and our Azure Synapse SQL Pool were created in an earlier experiment

```
Invoke-Sqlcmd -ServerInstance "adesqlserverxx.database.windows.net" -U sqladmin -P "Sql@Server@1234" -Database adesqldwxx -InputFile C:\projects\AzureDataEngineering\repos\azure-data-engineering\labs\Emp.sql
```

2. Find your SQL Server instance that was created in our previous experiment.



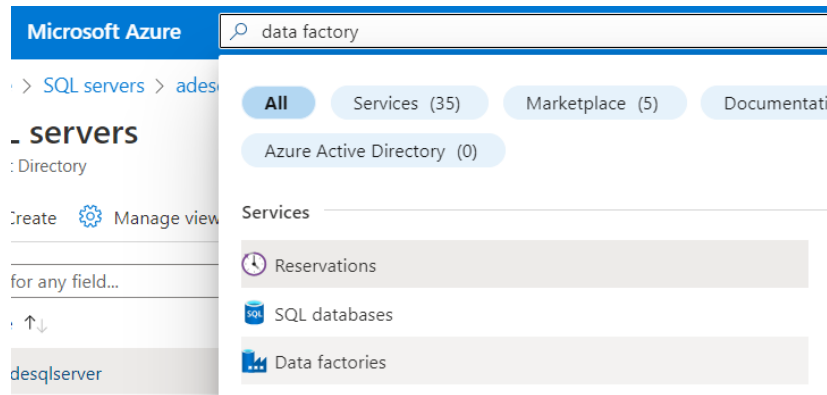
3. Allow Azure services to access for our SQL Server **adesqlserver** instance. Ensure that **Allow access to Azure services** is turned **ON** for your SQL Server so that Data Factory can write data to your SQL Server. To verify and turn on this setting, go to logical SQL server > Overview > Set server firewall> set the **Allow access to Azure services** option to **ON**.




Create a data factory

In this step, you create a data factory and start the Data Factory UI to create a pipeline in the data factory.

1. Open **Microsoft Edge** or **Google Chrome**. Currently, Data Factory UI is supported only in Microsoft Edge and Google Chrome web browsers. Use the search dialog to open the **Data Factories** blade.



2. Select **Create data factory**.



3. On the **Create Data Factory** page, under **Basics** tab, our default Azure **Subscription** should be select.
4. For **Resource Group**, take one of the following steps:
 - a. Select our existing **AzureDataEngineeringxx** resource group from the drop-down list.
5. Under **Region**, select the **East US 2** location for the data factory. Only locations that are supported are displayed in the drop-down list. The data stores (for example, Azure Storage and SQL Database) and computes (for example, Azure HDInsight) used by the data factory can be in other regions. We'll use **East US 2**.
6. Under **Name**, enter **ADFExperimentDataFactoryxx**. Remember to substitute your student number for the **xx**, for example **student1** would use **ADFExperimentDataFactory01**

The name of the Azure data factory must be *globally unique*. If you receive an error message about the name value, enter a different name for the data factory. (for example, ADFExperimentDataFactory01). For naming rules for Data Factory artifacts, see [Data Factory naming rules](#).

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ ✓

Resource group * ⓘ ✓

[Create new](#)

Instance details

Name * ⓘ ✓

Region * ⓘ ✓

Version * ⓘ ✓

7. Under **Version**, select **V2**.
8. Select **Git configuration** tab on the top, and select the **Configure Git later** check box.

Create Data Factory ...

Basics Git configuration Networking Advanced

Azure Data Factory allows you to configure a Git repository system that allows for easier change tracking and collaboration. [Learn more about Git integration in Azure Data Factory](#)

Configure Git later ⓘ ☒

9. Select **Review + create**

✓ Validation Passed

TERMS

By clicking "Create", I (a) agree to the legal terms and privacy statement(s) associated with the Marketplace offering(s) listed above; (b) authorize Microsoft to bill my current payment method for the fees associated with the offering(s), with the same billing frequency as my Azure subscription; and (c) agree that Microsoft may share my contact, usage and transactional information with the provider(s) of the offering(s) for support, billing and other transactional activities. Microsoft does not provide rights for third-party offerings. See the [Azure Marketplace Terms](#) for additional details.

Basics

Subscription	Azure subscription 1
Resource group	AzureDataEngineering
Name	ADFExperimentDataFactoryxx01
Region	East US 2
Version	V2 (Recommended)

Networking

Connect via	Public endpoint
-------------	-----------------




Create

< Previous





Next


[Download a template for automation](#)


10. Select **Create** after the validation is passed.


 **Microsoft.DataFactory-20220519072447** | Overview  


Deployment


<<  Delete  Cancel  Redeploy  Refresh


 Overview


 Inputs


 Outputs

 Template

 We'd love your feedback! →


 **Deployment is in progress**

 Deployment name: Microsoft.DataFactory-20220519072447
Subscription: [Azure subscription 1](#)
Resource group: [AzureDataEngineering](#)

 Deployment details [\(Download\)](#)

11. After the creation is finished, you see the notice in Notifications center. Select **Go to resource** to navigate to the Data factory page.

✓ Your deployment is complete

 Deployment name: Microsoft.DataFactory-2022051907...
Subscription: [Azure subscription 1](#)
Resource group: [AzureDataEngineering](#)


✓ Deployment details [\(Download\)](#)

✓ Next steps

[Go to resource](#)

12. Select **Open** on the **Open Azure Data Factory Studio** tile to launch the Azure Data Factory UI in a separate tab.

Getting started

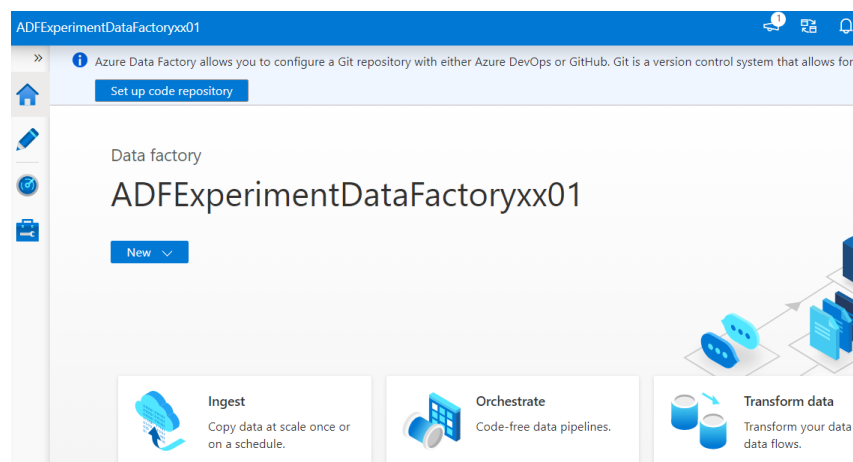


Create a pipeline

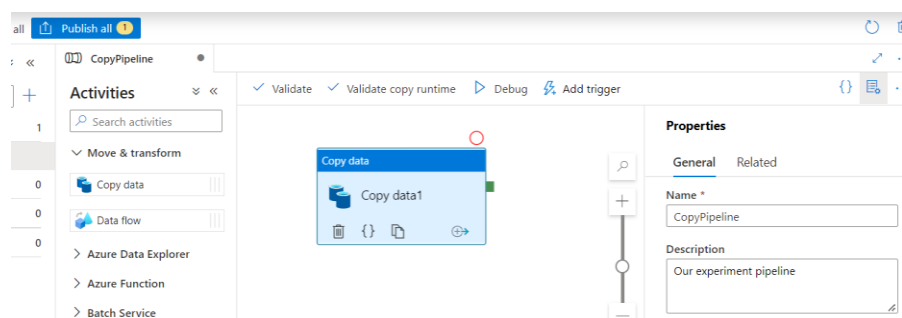
In this step, you create a pipeline with a copy activity in the data factory.

The copy activity copies data from Blob storage to SQL Database.

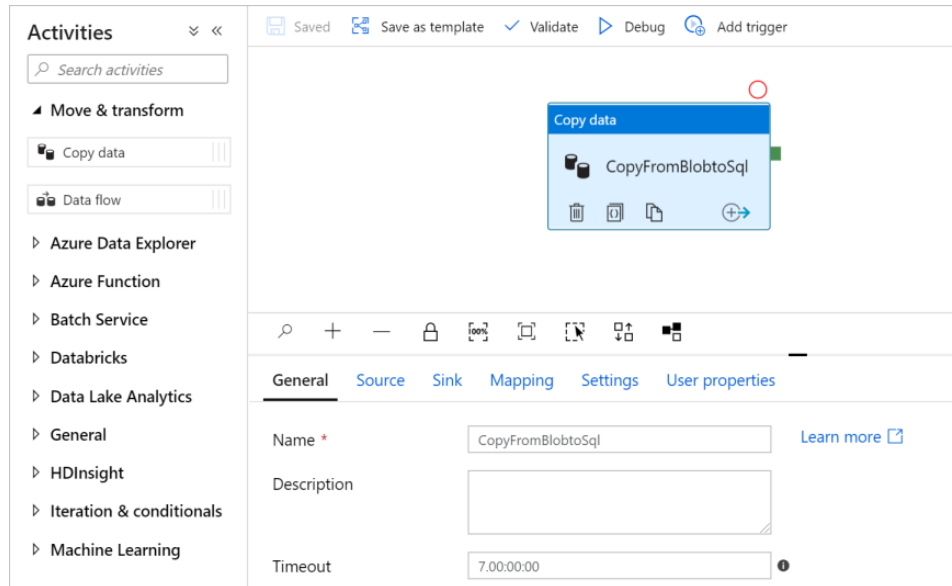
1. On the home page, select **Orchestrate**.



1. In the General panel under **Properties**, specify **CopyPipeline** for **Name**. Then collapse the panel by clicking the Properties icon in the top-right corner.

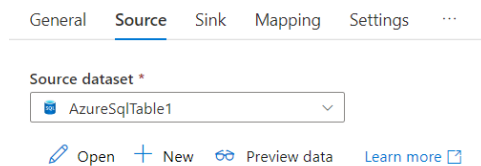


2. In the **Activities** tool box, expand the **Move and Transform** category, and drag and drop the **Copy Data** activity from the tool box to the pipeline designer surface. Specify **CopyFromBlobToSql** for **Name**.

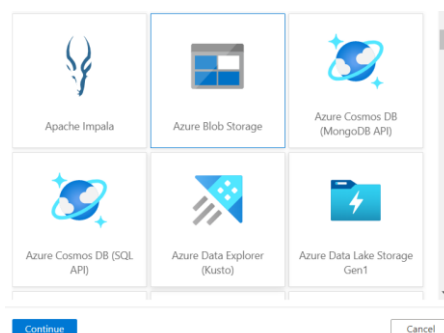


Configure source

1. Go to the **Source** tab. Select **+ New** to create a source dataset.




2. In the **New Dataset** dialog box, select **Azure Blob Storage**, and then select **Continue**. The source data is in Blob storage, so you select **Azure Blob Storage** for the source dataset.




3. In the **Select Format** dialog box, choose the format type of your data, and then select **Continue**.

Select format


Choose the format type of your data



Avro



Binary



DelimitedText

Continue Back Cancel

- In the **Set Properties** dialog box, enter **SourceBlobDataset** for Name. Select the checkbox for **First row as header**. Under the **Linked service** text box, select **+ New**.

New linked service

Azure Blob Storage [Learn more](#)

Account key

Connection string Azure Key Vault

Account selection method

☒ From Azure subscription ☐ Enter manually

Azure subscription

Azure subscription 1 (b7e892b5-020c-44c1-8d40-99e608f1b2d6)

Storage account name *

adestoragepowershell

Create Cancel Test connection

- In the **New Linked Service (Azure Blob Storage)** dialog box, enter **AzureStorageLinkedService** as name, select your storage account from the **Storage account name** list.

Connection successful

Create Cancel Test connection

- Test connection, select **Create** to deploy the linked service.
- After the linked service is created, it's navigated back to the **Set properties** page. Next to **File path**, select the folder icon next to the folder/file input boxes.

Set properties

Name

SourceBlobDataset

Linked service *

AzureBlobStorage1

File path

input / Directory / emp.txt

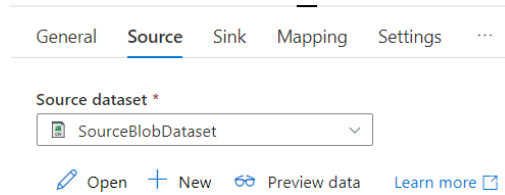
First row as header ☒

Import schema

☒ From connection/store ☐ From sample file ☐ None

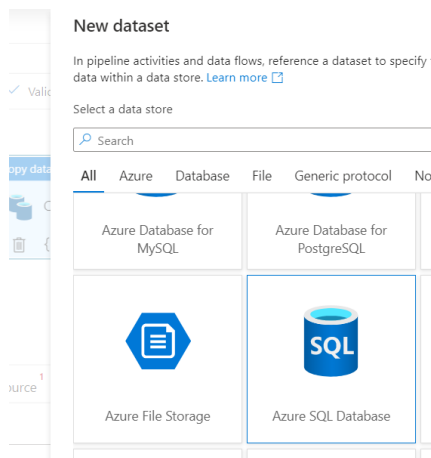
OK Back Cancel

8. Navigate to the **input** folder, select the **emp.txt** file, and then select **OK**.
9. Select **OK**. It automatically navigates to the pipeline page. In **Source** tab, confirm that **SourceBlobDataset** is selected. To preview data on this page, select **Preview data**.




Configure sink

1. Go to the **Sink** tab, and select **+ New** to create a sink dataset.
2. In the **New Dataset** dialog box, input "SQL" in the search box to filter the connectors, select **Azure SQL Database**, and then select **Continue**. In this experiment, you copy data to a SQL database.





3. In the **Set Properties** dialog box, enter **OutputSqlDataset** for Name. From the **Linked service** dropdown list, select **+ New**.




4. A dataset must be associated with a linked service. The linked service has the connection string that Data Factory uses to connect to SQL Database at runtime. The dataset specifies the container, folder, and the file (optional) to which the data is copied.
5. In the **New Linked Service (Azure SQL Database)** dialog box, take the following steps:

- a. Under **Name**, enter **AzureSqlDatabaseLinkedService**.
- b. Under **Server name**, select your SQL Server instance.
- c. Under **Database name**, select your **adesqldwxx**.
- d. Under **User name**, enter the name of the user.
- e. Under **Password**, enter the password for the user.
- f. Select **Test connection** to test the connection.
- g. Select **Create** to deploy the linked service.

Edit linked service

 Azure SQL Database [Learn more](#) 

Account selection method 

☐ From Azure subscription ☒ Enter manually

Fully qualified domain name *

adesqlserver.database.windows.net

Database name *

adesqldw

Authentication type *

SQL authentication

User name *


sqladmin

Password

Azure Key Vault

Password *

.....

Always encrypted 

☐

Additional connection properties

6. It automatically navigates to the **Set Properties** dialog box. In **Table**, select **[dbo].[emp]**. Then select **OK**.

Name
AzureSqlTable1

Linked service *
AzureSqlDatabase1

Table name
dbo.emp

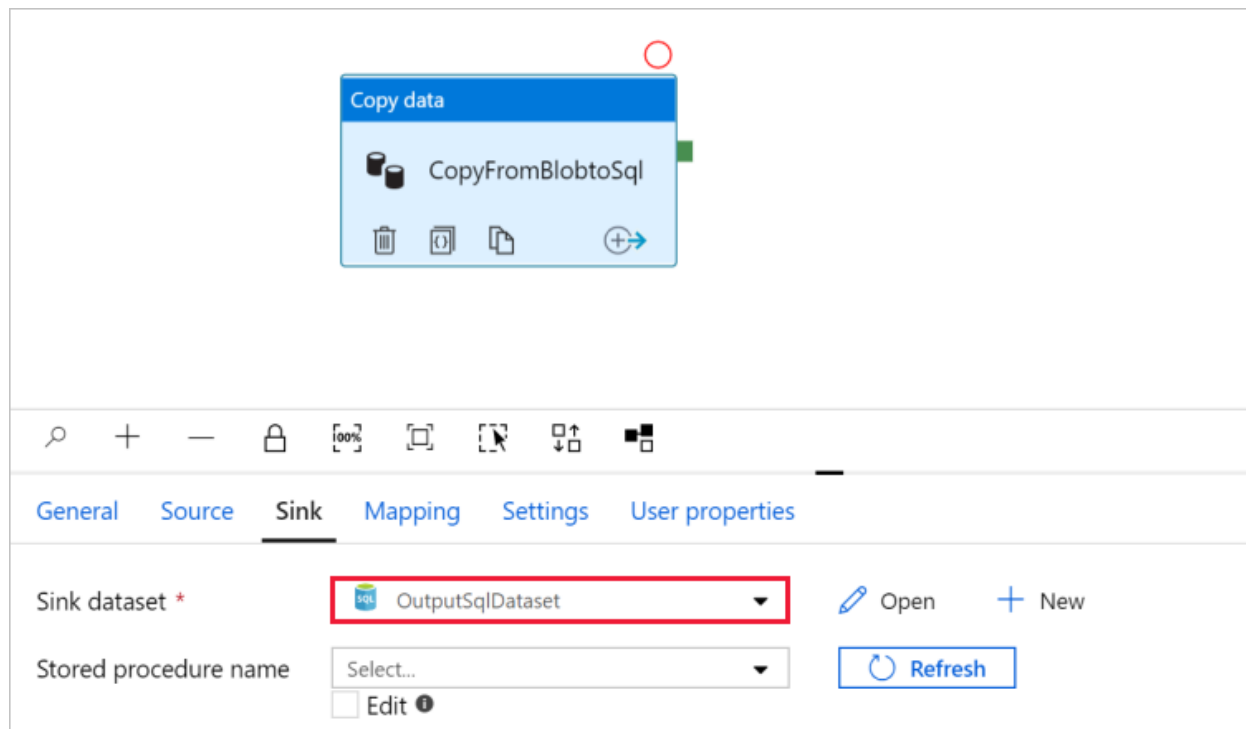
☐ Edit

Import schema
☒ From connection/store
 ☐ None

> Advanced

OK Back


- Go to the tab with the pipeline, and in **Sink Dataset**, confirm that **OutputSqlDataset** is selected.



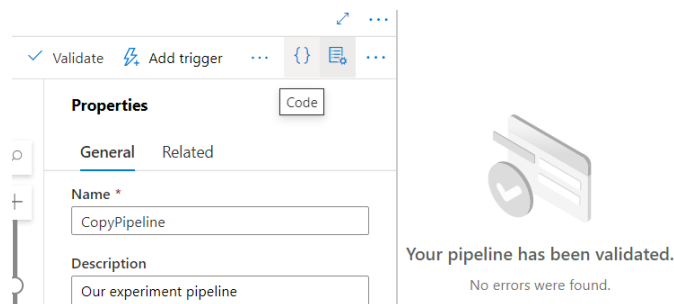
You can optionally map the schema of the source to corresponding schema of destination by following [Schema mapping in copy activity](#).

Validate the pipeline

To validate the pipeline, select **Validate** from the tool bar.




You can see the JSON code associated with the pipeline by clicking **Code** on the upper right.



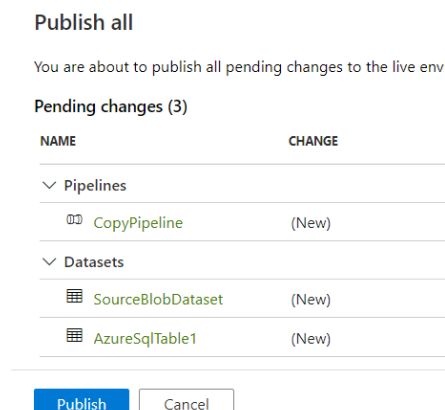
Debug and publish the pipeline

You can debug a pipeline before you publish artifacts (linked services, datasets, and pipeline) to Data Factory or your own Azure Repos Git repository.

1. To debug the pipeline, select **Debug** on the toolbar. You see the status of the pipeline run in the **Output** tab at the bottom of the window.
2. Once the pipeline can run successfully, in the top toolbar, select **Publish all**. This action publishes entities (datasets, and pipelines) you created to Data Factory.



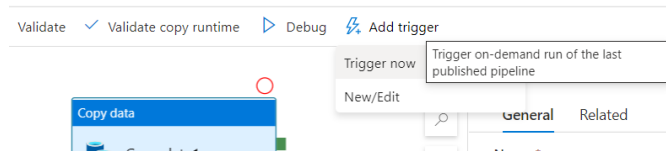
3. Wait until you see the **Successfully published** message. To see notification messages, click the **Show Notifications** on the top-right (bell button).



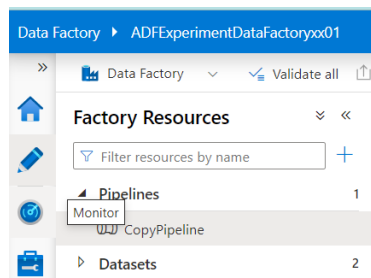
Trigger the pipeline manually

In this step, you manually trigger the pipeline you published in the previous step.

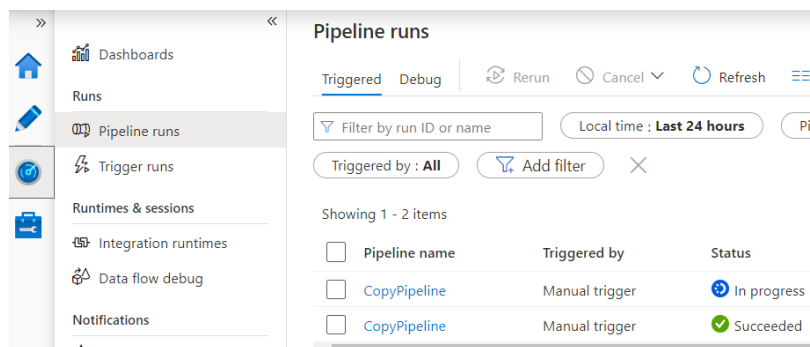
1. Select **Trigger** on the toolbar, and then select **Trigger Now**. On the **Pipeline Run** page, select **OK**.



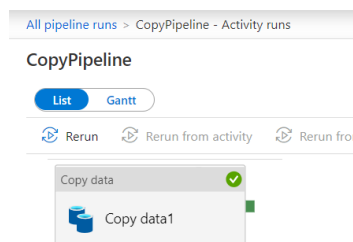
2. Go to the **Monitor** tab on the left.



3. You see a pipeline run that is triggered by a manual trigger. You can use links under the **PIPELINE NAME** column to view activity details and to rerun the pipeline.



4. To see activity runs associated with the pipeline run, select the **CopyPipeline** link under the **PIPELINE NAME** column. In this example, there's only one activity, so you see only one entry in the list. For details about the copy operation, select the **Details** link (eyeglasses icon) under the **ACTIVITY NAME** column.



5. Select **All pipeline runs** at the top to go back to the Pipeline Runs view. To refresh the view, select **Refresh**.
6. Verify that two more rows are added to the **emp** table in the database.

```
Invoke-Sqlcmd -ServerInstance "adesqlserverxx.database.windows.net" -U  
sqladmin -P "Sql@Server@1234" -Database adesqldwxx -Query "select *  
from emp"
```

Note: Substitute the correct naming for your ServerInstance endpoint URL and database.

```
PS C:\projects\AzureDataEngineering> Invoke-Sqlcmd -ServerInstance "adesqlserver  
"select * from emp"  
  
ID FirstName LastName  
--  
22 John Doe  
60 John Doe  
82 Jane Doe  
120 Jane Doe  
  
PS C:\projects\AzureDataEngineering>
```

Trigger the pipeline on a schedule

In this schedule, you create a schedule trigger for the pipeline. The trigger runs the pipeline on the specified schedule, such as hourly or daily. Here you set the trigger to run every minute until the specified end datetime.

1. Go to the **Author** tab on the left above the monitor tab.
2. Go to your pipeline, click **Trigger** on the tool bar, and select **New/Edit**.
3. In the **Add triggers** dialog box, select **+ New** for **Choose trigger** area.
4. In the **New Trigger** window, take the following steps:
 - a. Under **Name**, enter **RunEveryMinute**.
 - b. Update the **Start date** for your trigger. If the date is before current datetime, the trigger will start to take effect once the change is published.
 - c. Under **Time zone**, select the drop-down list.
 - d. Set the **Recurrence** to **Every 1 Minute(s)**.
 - e. Select the checkbox for **Specify an end date**, and update the **End On** part to be a few minutes past the current datetime. The trigger is activated only after you publish the changes. If you set it to only a couple of minutes apart, and you don't publish it by then, you don't see a trigger run.
 - f. For **Activated** option, select **Yes**.

g. Select **OK**.

Important

A cost is associated with each pipeline run, so set the end date appropriately.

5. On the **Edit trigger** page, review the warning, and then select **Save**. The pipeline in this example doesn't take any parameters.
6. Click **Publish all** to publish the change.
7. Go to the **Monitor** tab on the left to see the triggered pipeline runs.

Pipeline runs					
Time : Last 24 hours (4/12/20 4:34 PM ...)		Time zone : Beijing, Chongqing, Hong Kong...		Runs : Latest runs	
All status ▾		Rerun	Cancel ▾	Refresh	Edit columns
Showing 1 - 4 items					
<input type="checkbox"/> PIPELINE NAME	RUN START ↑↓	DURATION	TRIGGERED BY	STATUS	PARAMETERS
<input type="checkbox"/> CopyPipeline	4/13/20, 4:34:00 PM	00:00:09	RunEveryMinute	✓ Succeeded	
<input type="checkbox"/> CopyPipeline	4/13/20, 4:32:59 PM	00:00:11	RunEveryMinute	✓ Succeeded	
<input type="checkbox"/> CopyPipeline	4/13/20, 4:31:59 PM	00:00:11	RunEveryMinute	✓ Succeeded	

8. To switch from the **Pipeline Runs** view to the **Trigger Runs** view, select **Trigger Runs** on the left side of the window.
9. You see the trigger runs in a list.
10. Verify that two rows per minute (for each pipeline run) are inserted into the **emp** table until the specified end time.